



Reading accuracy and complete-pipeline cost

Frozen models: six-condition accuracy and batch-1 FP32 timing

(a) Error and latency



(b) Pipeline parameters



(c) Peak CUDA memory

